

# METCS777-term-project-code-sample-EDA-Team17

December 5, 2025

```
[1]: from pyspark.sql import SparkSession
from pyspark.sql.functions import col, count, isnan, when, expr
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

[2]: spark = SparkSession.builder \
    .appName("Airbnb_EDA") \
    .getOrCreate()

df_la = spark.read.option("header", True).csv("/Users/yibingwang/Desktop/777/
˓term project/la_cleaned.csv")
df_ny = spark.read.option("header", True).csv("/Users/yibingwang/Desktop/777/
˓term project/ny_cleaned.csv")

WARNING: Using incubator modules: jdk.incubator.vector
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).
25/12/05 20:05:20 WARN NativeCodeLoader: Unable to load native-hadoop library
for your platform... using builtin-java classes where applicable

[3]: # Missing value summary LA vs NY
def missing_table(df):
    return df.select([(count(when(col(c).isNull(), c)).alias(c)) for c in df.
˓columns]).toPandas()

print("Missing - LA")
print(missing_table(df_la))

print("\nMissing - NY")
print(missing_table(df_ny))
```

Missing - LA

	id	neighbourhood_group	neighbourhood	latitude	longitude	room_type	\
0	0	0	0	0	0	0	0

```

    price   minimum_nights   number_of_reviews   last_review   reviews_per_month   \
0          0                  0                      0           199                   0

    calculated_host_listings_count   availability_365   number_of_reviews_ltm   \
0                           0                      0                   0

    license
0        534

Missing - NY
    id   neighbourhood_group   neighbourhood   latitude   longitude   room_type   \
0     0                  0                      0           1           1           0

    price   minimum_nights   number_of_reviews   last_review   reviews_per_month   \
0          0                  1                      0           175                   0

    calculated_host_listings_count   availability_365   number_of_reviews_ltm   \
0                           1                      0                   0

    license
0        453

```

```
[4]: #Descriptive statistics (key metrics comparison: count, min, mean, max, 25%, 50%, 75%)
#The goal is to compare LA vs NY listing structure - which city is more expensive,
# where prices are more concentrated, and where we see more extreme values.
desc_la = df_la.describe().toPandas()
desc_ny = df_ny.describe().toPandas()

print("LA Summary Stats:")
print(desc_la)

print("\nNY Summary Stats:")
print(desc_ny)
```

25/12/05 20:05:23 WARN SparkStringUtils: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

LA Summary Stats:

	summary	id	neighbourhood_group	neighbourhood
0	count	779	779	779
1	mean	8.012098321717916E17	None	None
2	stddev	5.614733099566928E17	None	None
3	min	1001638692383599358	City of Los Angeles	Acton
4	max	998166923867473626	Unincorporated Areas	Woodland Hills

	latitude	longitude	room_type	\
0	779	779	779	
1	34.05233484813748	-118.31044681476178	None	
2	0.1541243801756928	0.17906876621250054	None	
3	33.3403891	-117.71752	Entire home/apt	
4	34.70167922973633	-118.90249	Shared room	

	price	minimum_nights	number_of_reviews	last_review	\
0	779	779	779	580	
1	400.7586649550706	18.319640564826702	37.094993581514764	None	
2	2076.5938545604454	26.005422845079067	73.35939080390115	None	
3	100.0	1.0	0.0	2014-06-15	
4	99.0	90.0	98.0	2025-09-01	

	reviews_per_month	calculated_host_listings_count	availability_365	\
0	779	779	779	
1	1.2037997432605898	21.50834403080873	260.9653401797176	
2	1.6170927262515171	68.80700904173692	108.93195008035795	
3	0.0	1.0	0.0	
4	9.85	91.0	98.0	

	number_of_reviews_ltm	license		
0	779	245		
1	9.397946084724005	180267.25		
2	16.573922462764106	84426.86084374198		
3	0.0	000000		
4	94.0	nrp21-00296		

NY Summary Stats:

	summary	id	neighbourhood_group	neighbourhood	\
0	count	570	570	570	
1	mean	6.421791455434981E17	4.94585533E8	None	
2	stddev	5.3284011391513235E17	None	None	
3	min	1002823881758639223	494585533	Allerton	
4	max	993514	Staten Island	Woodside	

	latitude	longitude	room_type	price	\
0	569	569	570	570	
1	40.72599369336432	-73.94593680618843	40.5793	873.9877495263158	
2	0.0582334657170787	0.05581482270314689	None	5265.964833184427	
3	40.56807	-73.72869	40.5793	-73.98277	
4	40.9031966657953	-74.13372	Shared room	995.0	

	minimum_nights	number_of_reviews	last_review	reviews_per_month	\
0	569	570	395	570	
1	26.047451669595784	33.14912280701754	None	0.7718245614035084	
2	18.536399231985065	64.04252657512423	None	1.5701408164806276	
3	1.0	0.0	2016-03-19	0.0	

```

4          90.0        99.0  2025-10-01        9.59
calculated_host_listings_count    availability_365 number_of_reviews_ltm \
0                      569                  570                  570
1      93.68365553602811  248.9149649122807  5.464912280701754
2     265.70487923368165  107.46275947298612 16.30453796500832
3                     1.0                  0.0                  0.0
4                     93.0                 99.0                 99.0

license
0          117
1         364.0
2           None
3          364
4  ose-strreg-0002161

```

```

[5]: #Price distribution Histogram + Log
#Convert to Pandas (select price only)
la_price = df_la.select("price").dropna().toPandas()
ny_price = df_ny.select("price").dropna().toPandas()

plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plt.hist(la_price["price"], bins=50)
plt.title("LA Price Distribution")

plt.subplot(1,2,2)
plt.hist(ny_price["price"], bins=50)
plt.title("NY Price Distribution")

plt.show()

#Log-transformed price distribution

la_price["price"] = pd.to_numeric(la_price["price"], errors="coerce")
ny_price["price"] = pd.to_numeric(ny_price["price"], errors="coerce")

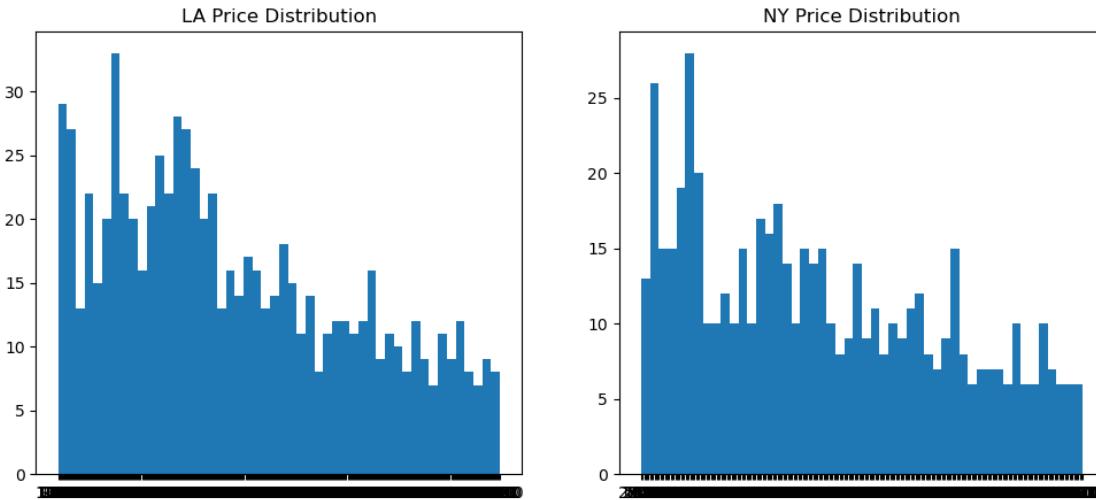
plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
plt.hist(np.log1p(la_price["price"].dropna()), bins=50)
plt.title("LA Log Price Distribution")

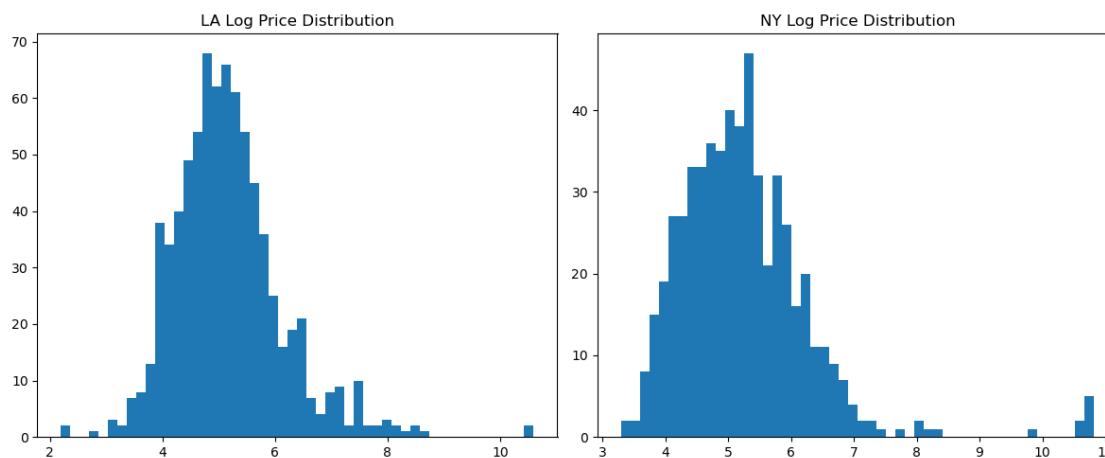
plt.subplot(1,2,2)
plt.hist(np.log1p(ny_price["price"].dropna()), bins=50)
plt.title("NY Log Price Distribution")

```

```
plt.tight_layout()  
plt.show()
```



```
/opt/anaconda3/lib/python3.12/site-packages/pandas/core/arraylike.py:399:  
RuntimeWarning: invalid value encountered in log1p  
    result = getattr(ufunc, method)(*inputs, **kwargs)
```



```
[6]: #Scatter Geo Latitude vs Longitude, color = price  
#Check price range and extreme values  
df_la.selectExpr(  
    "max(price)",  
    "percentile(price, 0.95)",  
    "percentile(price, 0.99)"
```

```

).show()

df_ny.selectExpr(
    "max(price)",
    "percentile(price, 0.95)",
    "percentile(price, 0.99")
).show()
#There are extreme price values observed in previous attempt
#Keep only rows with 0 < price <= 2000
#For safety, convert price to double explicitly first
df_la_num = df_la.withColumn(
    "price_double",
    expr("try_cast(price as double)")
)

df_ny_num = df_ny.withColumn(
    "price_double",
    expr("try_cast(price as double)")
)

df_la_cap = df_la_num.filter(
    (col("price_double") > 0.0) & (col("price_double") <= 2000.0)
)

df_ny_cap = df_ny_num.filter(
    (col("price_double") > 0.0) & (col("price_double") <= 2000.0)
)

df_la_cap.selectExpr("min(price_double) as min_price", "max(price_double) as max_price").show()
df_ny_cap.selectExpr("min(price_double) as min_price", "max(price_double) as max_price").show()

print("LA rows after cap:", df_la_cap.count())
print("NY rows after cap:", df_ny_cap.count())

#plot
la_geo = df_la_cap.select("latitude", "longitude", "price_double").dropna().toPandas()
ny_geo = df_ny_cap.select("latitude", "longitude", "price_double").dropna().toPandas()

print("LA geo shape:", la_geo.shape)
print("NY geo shape:", ny_geo.shape)
print(la_geo.dtypes)
print(la_geo.head())
# In la_geo, all three columns are object type in Pandas, so convert to numeric

```

```

la_geo["latitude"] = pd.to_numeric(la_geo["latitude"], errors="coerce")
la_geo["longitude"] = pd.to_numeric(la_geo["longitude"], errors="coerce")
la_geo["price_double"] = pd.to_numeric(la_geo["price_double"], errors="coerce")

ny_geo["latitude"] = pd.to_numeric(ny_geo["latitude"], errors="coerce")
ny_geo["longitude"] = pd.to_numeric(ny_geo["longitude"], errors="coerce")
ny_geo["price_double"] = pd.to_numeric(ny_geo["price_double"], errors="coerce")

plt.figure(figsize=(12,5))

plt.subplot(1,2,1)
sc1 = plt.scatter(la_geo["longitude"], la_geo["latitude"],
                  c=la_geo["price_double"], s=5)
plt.colorbar(sc1, label="Price")
plt.title("LA Geo Price Map (price 2000)")
plt.xlabel("longitude")
plt.ylabel("latitude")

plt.subplot(1,2,2)
sc2 = plt.scatter(ny_geo["longitude"], ny_geo["latitude"],
                  c=ny_geo["price_double"], s=5)
plt.colorbar(sc2, label="Price")
plt.title("NY Geo Price Map (price 2000)")
plt.xlabel("longitude")
plt.ylabel("latitude")

plt.tight_layout()
plt.show()

```

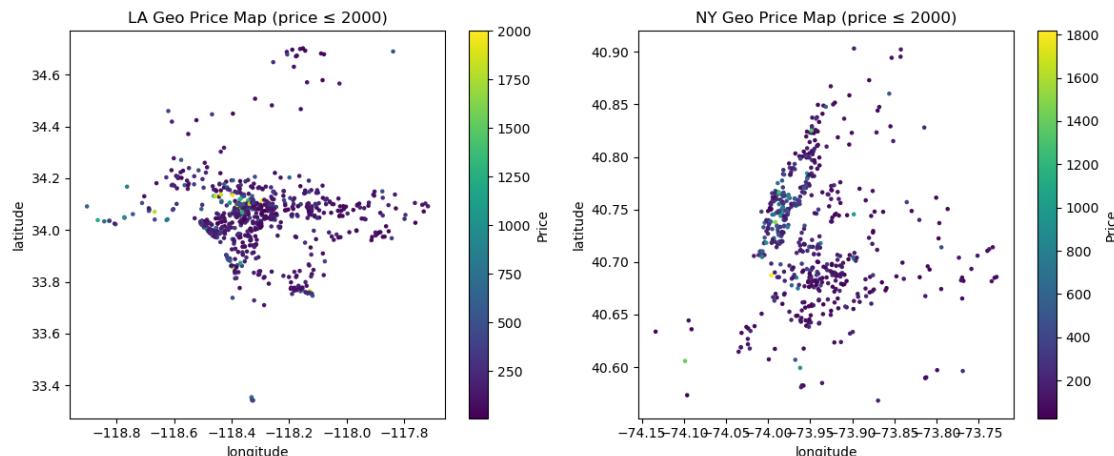
max(price) percentile(price, 0.95, 1) percentile(price, 0.99, 1)			
99.0  1015.699999999975  3232.140000000002			
max(price) percentile(price, 0.95, 1) percentile(price, 0.99, 1)			
995.0  871.349999999974  40000.0			
min_price max_price			
+-----+-----+			
8.0  2000.0			
+-----+-----+			

```

+-----+-----+
| min_price | max_price |
+-----+-----+
|     26.0 |   1817.0 |
+-----+-----+

LA rows after cap: 766
NY rows after cap: 556
LA geo shape: (766, 3)
NY geo shape: (556, 3)
latitude      object
longitude     object
price_double  float64
dtype: object
   latitude      longitude  price_double
0  34.16767360409868 -118.28108622078958      135.0
1          34.07628           -118.09548      114.0
2          33.84594           -118.11532      108.0
3          34.03572           -118.82613     938.0
4          34.0257            -118.36589      162.0

```



[7]: #Numeric features vs Price scatter plot (with log-transformed price)

```

la_pd = df_la.select("minimum_nights","price").dropna().toPandas()
ny_pd = df_ny.select("minimum_nights","price").dropna().toPandas()
for df in (la_pd, ny_pd):
    df["minimum_nights"] = pd.to_numeric(df["minimum_nights"], errors="coerce")
    df["price"] = pd.to_numeric(df["price"], errors="coerce")

plt.figure(figsize=(12,5))

```

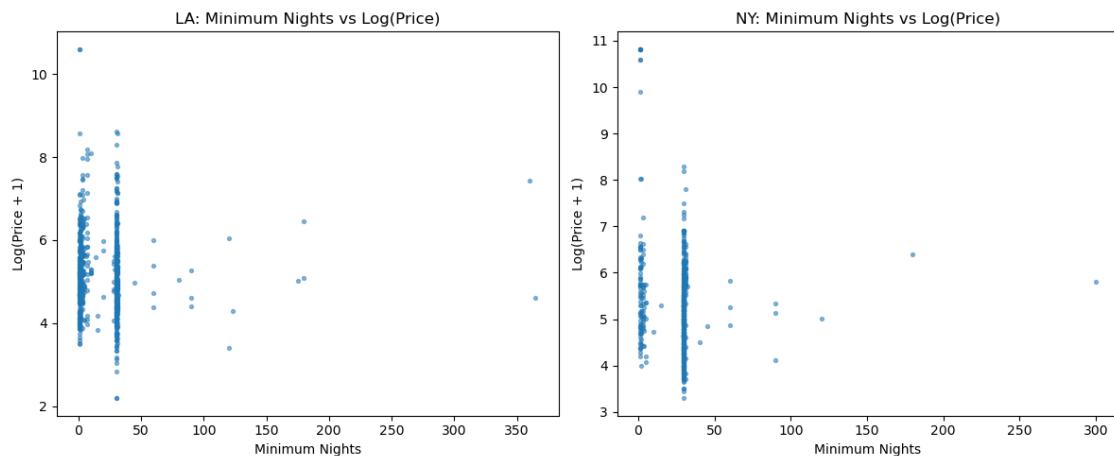
```

plt.subplot(1,2,1)
plt.scatter(la_pd["minimum_nights"] , np.log1p(la_pd["price"]),
           s=8, alpha=0.5)
plt.title("LA: Minimum Nights vs Log(Price)")
plt.xlabel("Minimum Nights")
plt.ylabel("Log(Price + 1)")

plt.subplot(1,2,2)
plt.scatter(ny_pd["minimum_nights"] , np.log1p(ny_pd["price"]),
           s=8, alpha=0.5)
plt.title("NY: Minimum Nights vs Log(Price)")
plt.xlabel("Minimum Nights")
plt.ylabel("Log(Price + 1)")

plt.tight_layout()
plt.show()

```



```

[8]: #Correlation Matrix Heatmap
numeric_cols = [
    "price",
    "minimum_nights",
    "number_of_reviews",
    "reviews_per_month",
    "availability_365",
    "calculated_host_listings_count",
    "number_of_reviews_ltm",
    "latitude",
    "longitude"
]

la_corr = df_la.select(*numeric_cols).toPandas().corr()

```

```

ny_corr = df_ny.select(*numeric_cols).toPandas().corr()

fig, axes = plt.subplots(1, 2, figsize=(20, 8))

#LA Heatmap
im1 = axes[0].imshow(la_corr, vmin=-1, vmax=1, cmap="coolwarm")
axes[0].set_title("LA Correlation Matrix", fontsize=18)
axes[0].set_xticks(range(len(numeric_cols)))
axes[0].set_yticks(range(len(numeric_cols)))
axes[0].set_xticklabels(numeric_cols, rotation=45, ha="right")
axes[0].set_yticklabels(numeric_cols)

cbar1 = fig.colorbar(im1, ax=axes[0], shrink=0.8)
cbar1.set_label("Correlation", fontsize=12)

#NY Heatmap
im2 = axes[1].imshow(ny_corr, vmin=-1, vmax=1, cmap="coolwarm")
axes[1].set_title("NY Correlation Matrix", fontsize=18)
axes[1].set_xticks(range(len(numeric_cols)))
axes[1].set_yticks(range(len(numeric_cols)))
axes[1].set_xticklabels(numeric_cols, rotation=45, ha="right")
axes[1].set_yticklabels(numeric_cols)

cbar2 = fig.colorbar(im2, ax=axes[1], shrink=0.8)
cbar2.set_label("Correlation", fontsize=12)

plt.tight_layout()
plt.show()

```

